

# INTRODUCTION TO CONCURRENCY

Macton@  
insomniacgames.com

Let's Start w/  
a well-known  
"problem" and  
work backward...

PROBLEM:

THERE IS NO LOCK-FREE  
VERSION OF A DOUBLY-  
LINKED LIST.

PROBLEM:

THERE IS NO LOCK-FREE  
VERSION OF A DOUBLY-  
LINKED LIST.

OF COURSE  
NOT!

VERSION OF A DOUBLY-  
LINKED LIST.

BUT  
WHY  
NOT?

LIST.

IT'S NOT  
A  
CONCURRENT  
DATA  
STRUCTURE!

PROBLEM:

THERE IS NO LOCK-FREE  
VERSION OF A DOUBLY-  
LINKED LIST.

IT'S NOT  
REALLY  
A PROBLEM.

PROBLEM:

THERE IS NO LOCK-FREE  
VERSION OF A DOUBLY-  
LINKED LIST.

MORE LIKE  
AN  
INTERESTING  
PUZZLE.

PROBLEM:

THERE IS NO LOCK-FREE  
VERSION OF A DOUBLY -  
LINKED LIST.

IT'S THE  
WRONG LEVEL  
OF ABSTRACTION  
FOR  
CONCURRENCY.

CREATING CONCURRENT  
DATA STRUCTURES  
REQUIRES AN EXTRA  
DIMENSION OF INFO

CREATING CONCURRENT  
DATA STRUCTURES  
REQUIRES AN EXTRA  
DIMENSION OF INFO

←  
THIS SEEMS  
OBVIOUS

DATA STRUCTURES  
REQUIRES AN EXTRA  
DIMENSION OF INFO

SOMETIMES,  
TRAN. DATA  
STRUCTS CAN  
BE USED,

←  
THIS SE  
OBVIO

DATA STRUCTURES  
REQUIRES AN EXTRA  
DIMENSION OF INFO

JUST LIKE  
SOMETIMES  
2D STRUCTS  
CAN BE USED  
IN 3D APPS.

←  
THIS S  
O BUI

# DIMENSION OF INFO

BUT ONLY  
IF YOU  
PRESUME  
CERTAIN  
THINGS.

# DIMENSION

Remember -  
It's always  
about the  
data!

BUT  
IF YOU  
PRESUME  
CERTAIN  
THINGS

Concurrency is  
a data problem,  
not a code problem.

Concurrency is  
a data problem,  
not a code problem.

↑

designing  
code-first  
will only  
overcomplicate

DOUBLY-LINKED LIST  
PRESUMES SEQUENTIAL  
ORDER.

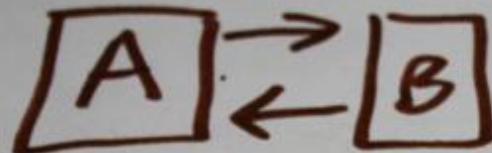
DOUBLY-LINKED LIST  
PRESUMES SEQUENTIAL  
ORDER.



Well,  
obviously it's  
a definition  
of an order, ...

RESUMES SEQUENTIAL  
ORDER.

But xforms  
of the data  
are also  
implicitly  
ordered.



e.g  
INSERT

- INSERT (C) AFTER (A)
- INSERT (D) AFTER (A)

HAS GUARANTEED RESULT:



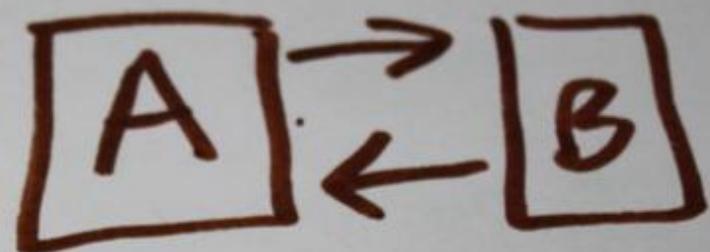
$\vdash [B]$

$\vdash (C)$   
 $\vdash (D)$

e.g.  
INSERT

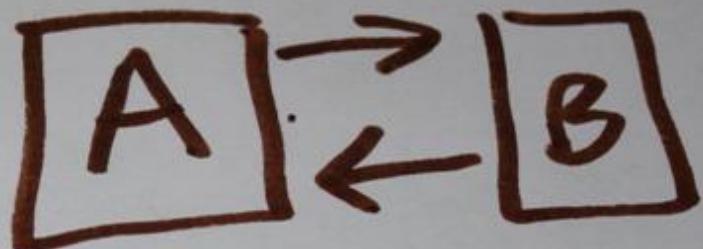
All ops  
have  
predictable  
results

The data  
struct only  
exists to  
facilitate  
the results.



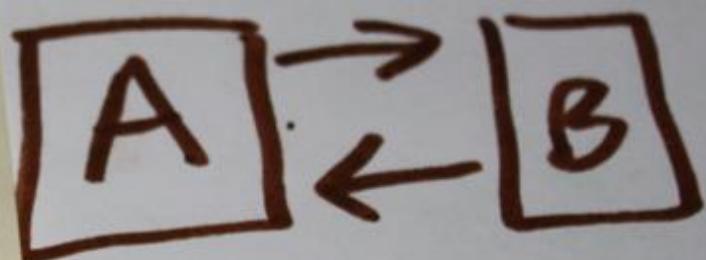
- INSERT (C)
- INSERT (D)

What  
Should The  
results  
be?



- INSERT (C)
- INSERT (D)

the  
problem  
isn't even  
the same



- INSERT (O)
- INSERT (Ω)

So what  
part of  
the solution  
would  
be the same?



- INS
- INS

Problem:  
Concurrent  
insert op

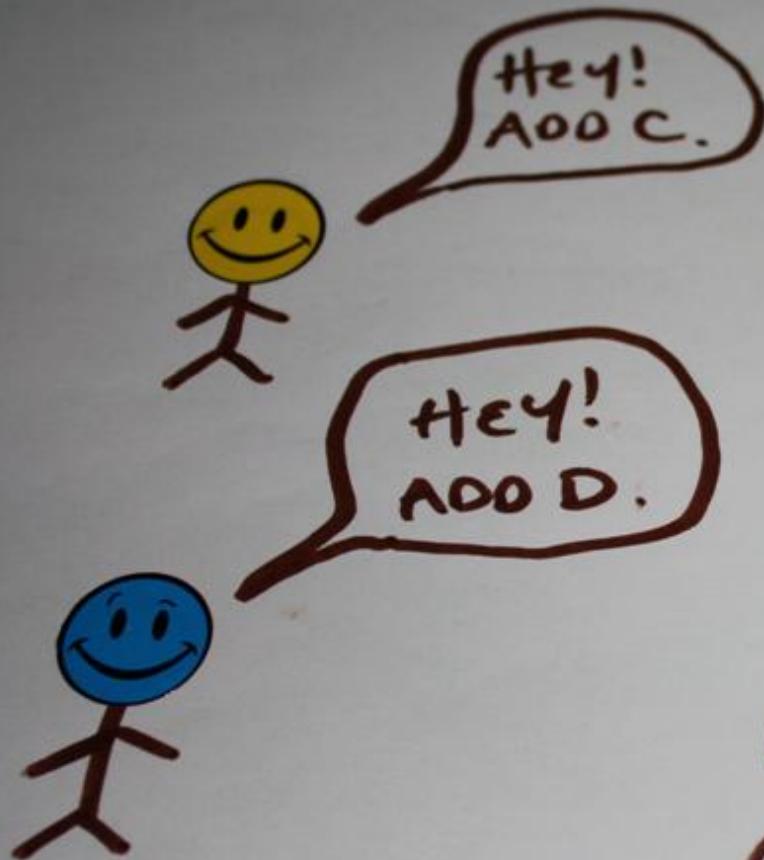
Problem:  
Concurrent  
insert of

What does  
it mean?

Problem:  
Concurrent  
insert of

T. i.e.

can't  
define data  
before even  
knowing the  
problem



WHAT CAN YOU  
GUARANTEE ABOUT  
THE ORDER HERE?



HEY!  
ADD C.



Don't ask,  
"which  
one was  
first?"



WHAT CAN YOU  
GUARANTEE ABOUT



HEY!  
ADD C.



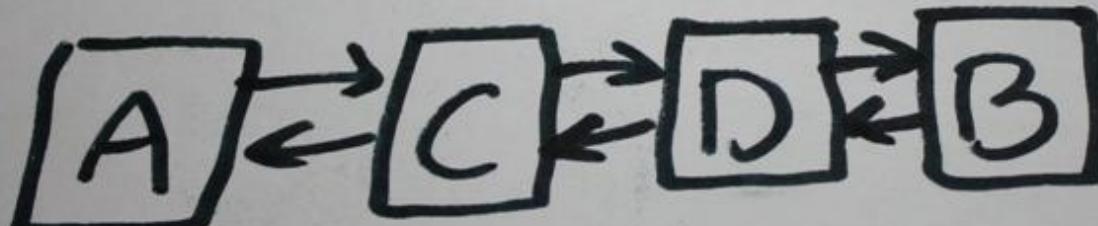
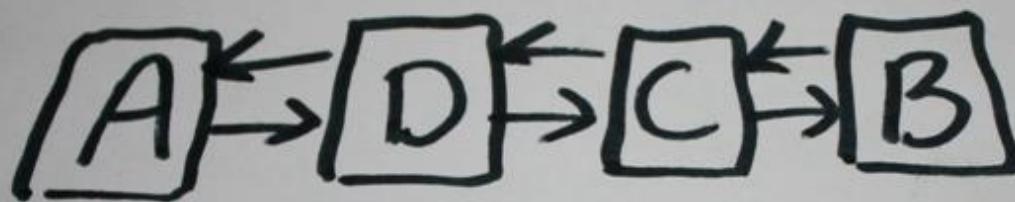
HEY!  
ADD D.

ASK, why  
does order  
matter for  
The problem.

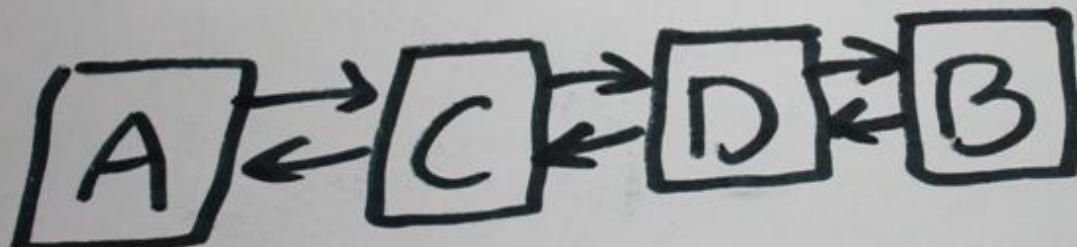
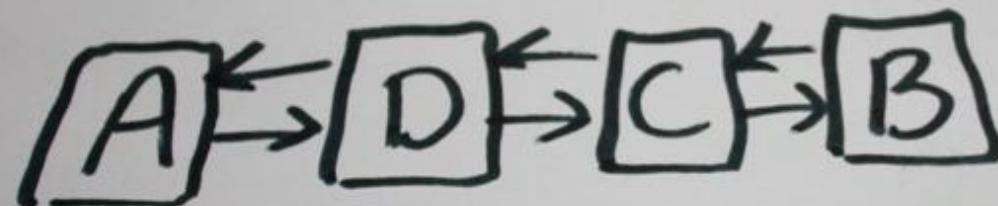


WHAT CAN YOU  
SUGGEST AS

WHICH ORDER IS "CORRECT" ?

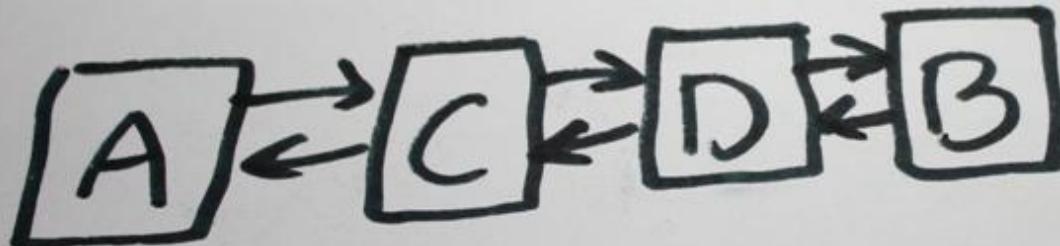


WHICH ORDER IS "CORRECT" ?



BOTH !

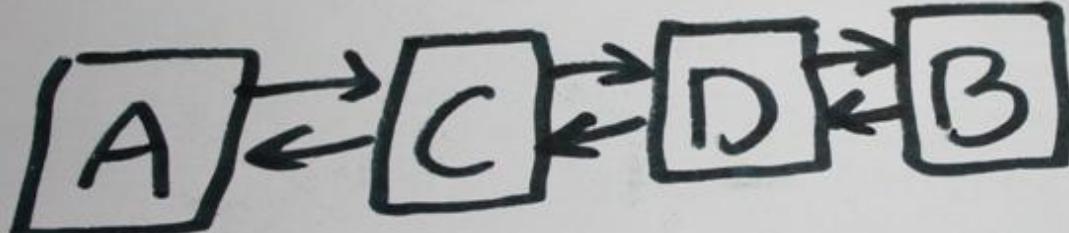
WHICH ORDER IS "CORRECT" ?



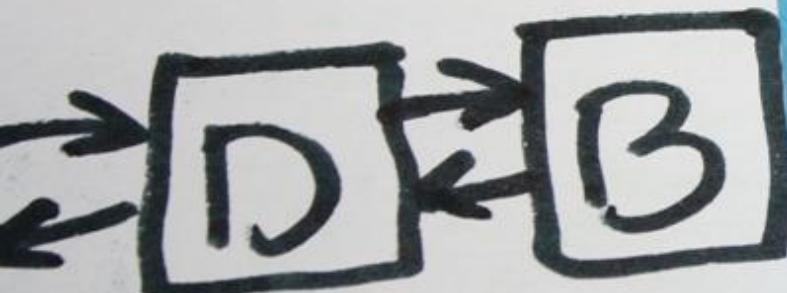
BOTH !

NEITHER !

WHICH ORDER IS "CORRECT"?



DEPENDS  
ON  
CONTEXT.



"CONTEXT"  
INCLUDES  
EXPLICIT  
ORDERING  
RULES.

---

BUT FIRST...

WHAT IS CONCURRENCY?

# WHAT IS CONCURRENCY?

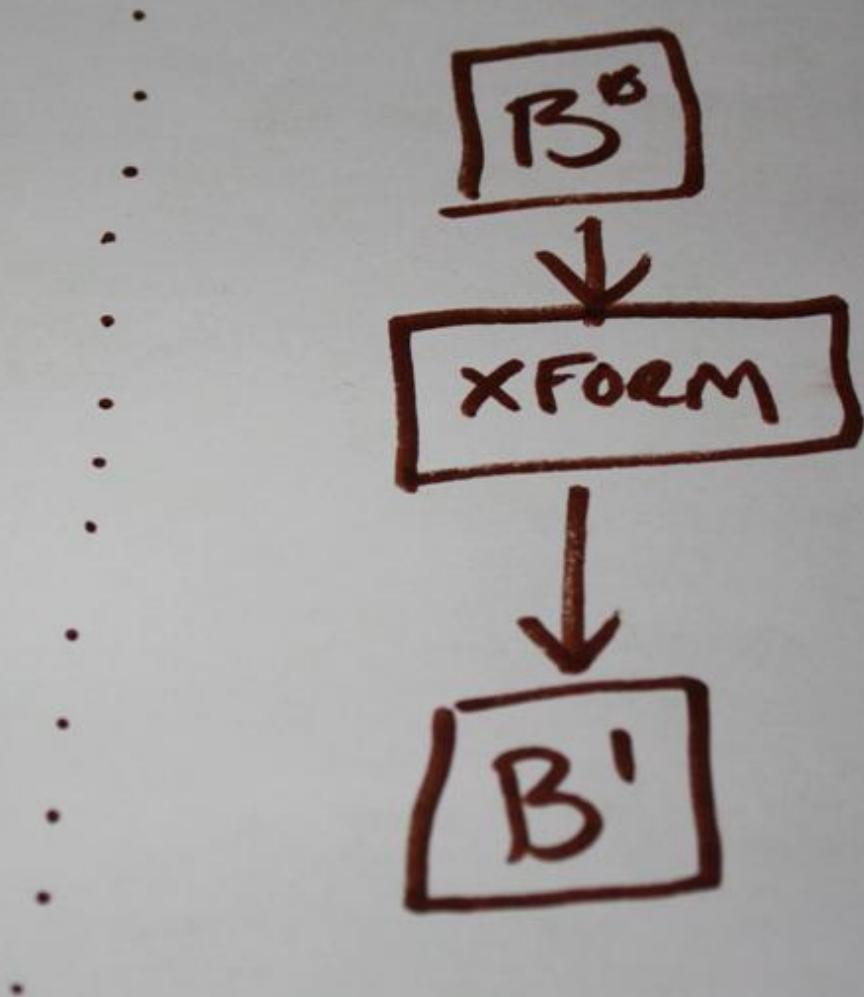
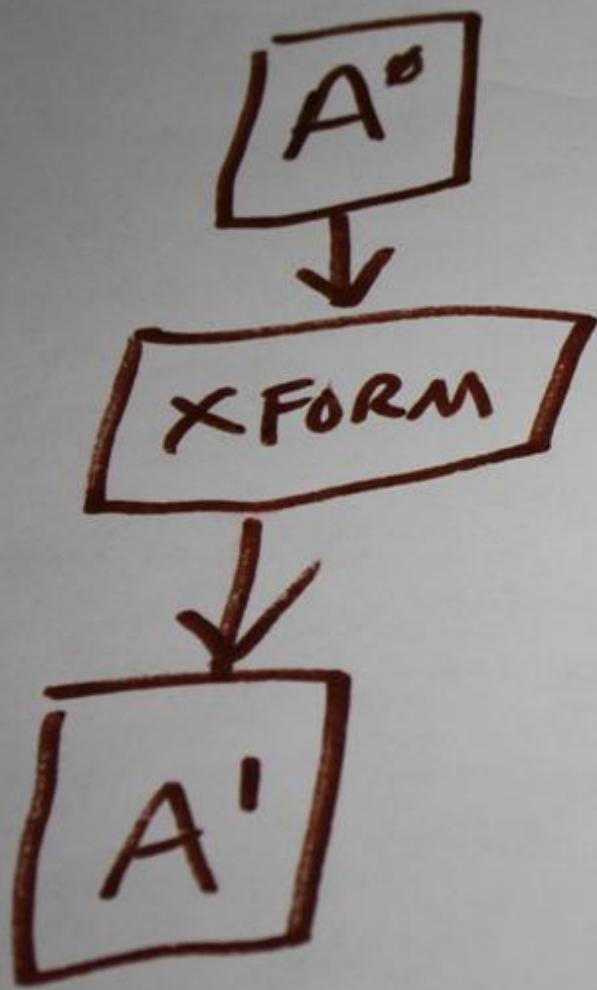
PARALLELISM  
VS.  
CONCURRENCY  
ARGUMENTS...

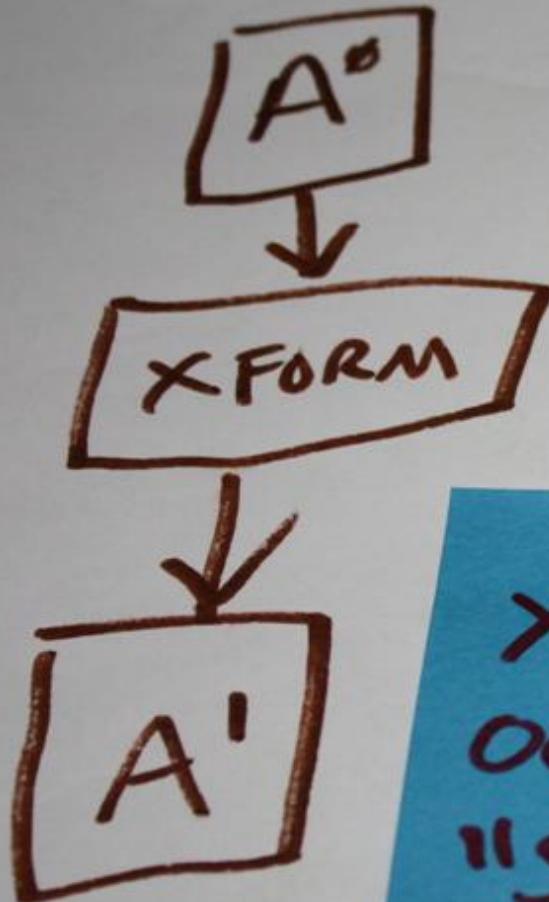
WHO CARES?!

NOT  
IMPORTANT.

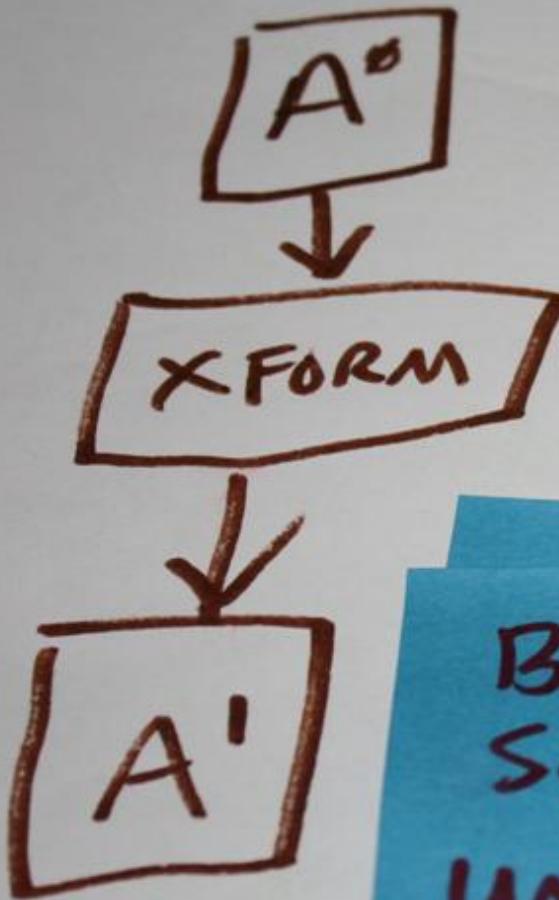
LEAVE IT.

CONCURRENCY IS  
TRANSFORMATION OF  
SHARED DATA SET.



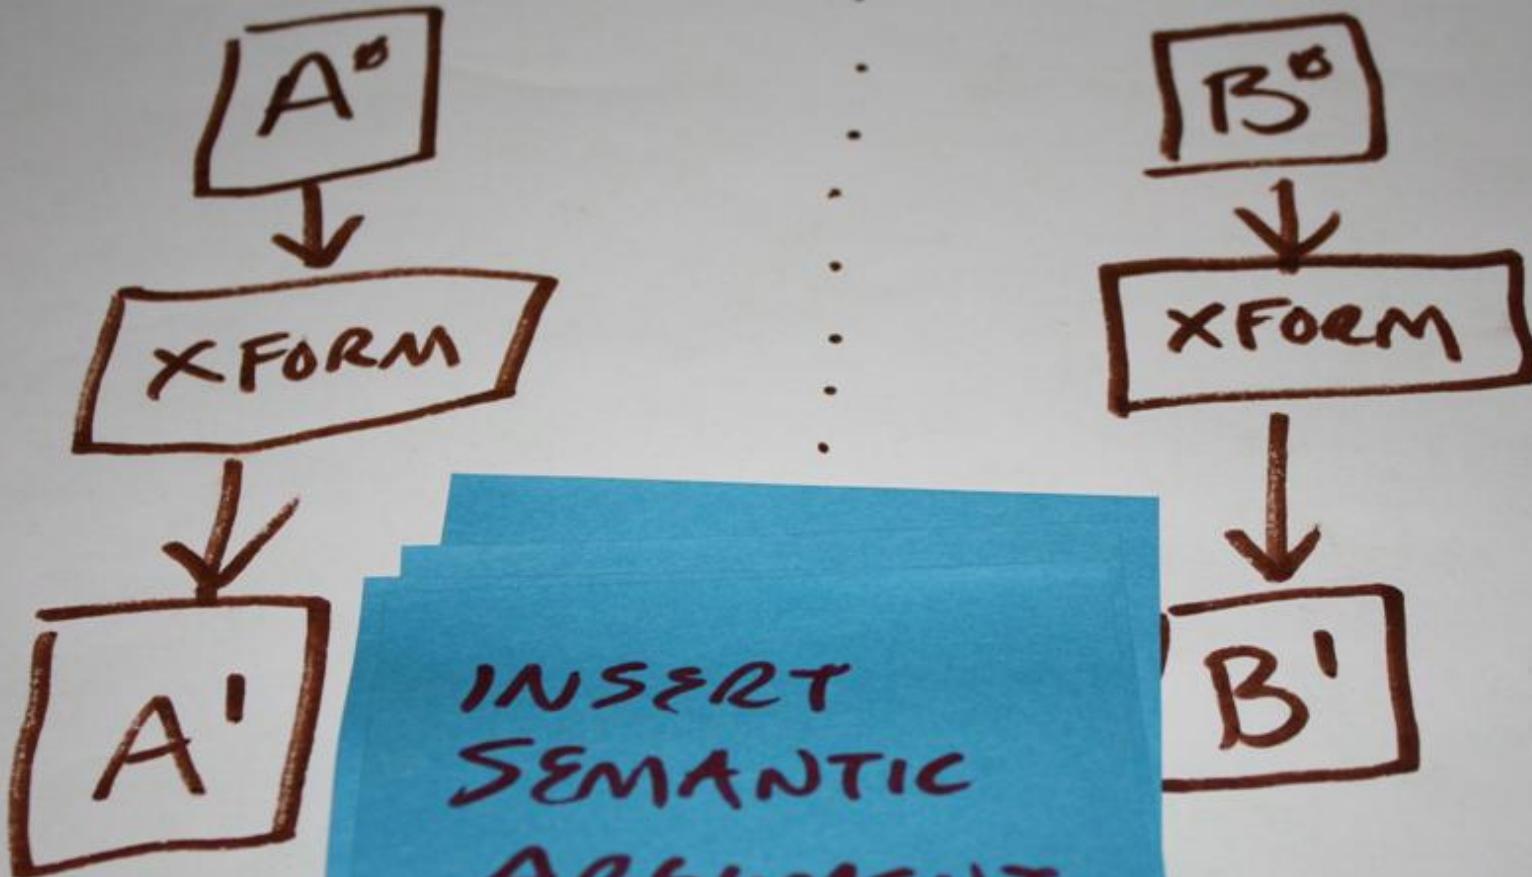


XFORMS  
OCCUR AT  
"SAME  
TIME"



BUT DATA  
SETS ARE  
UNRELATED.





INSERT  
SEMANTIC  
ARGUMENT  
HERE.

No  
concurrency  
...  
in any  
way that  
matters.

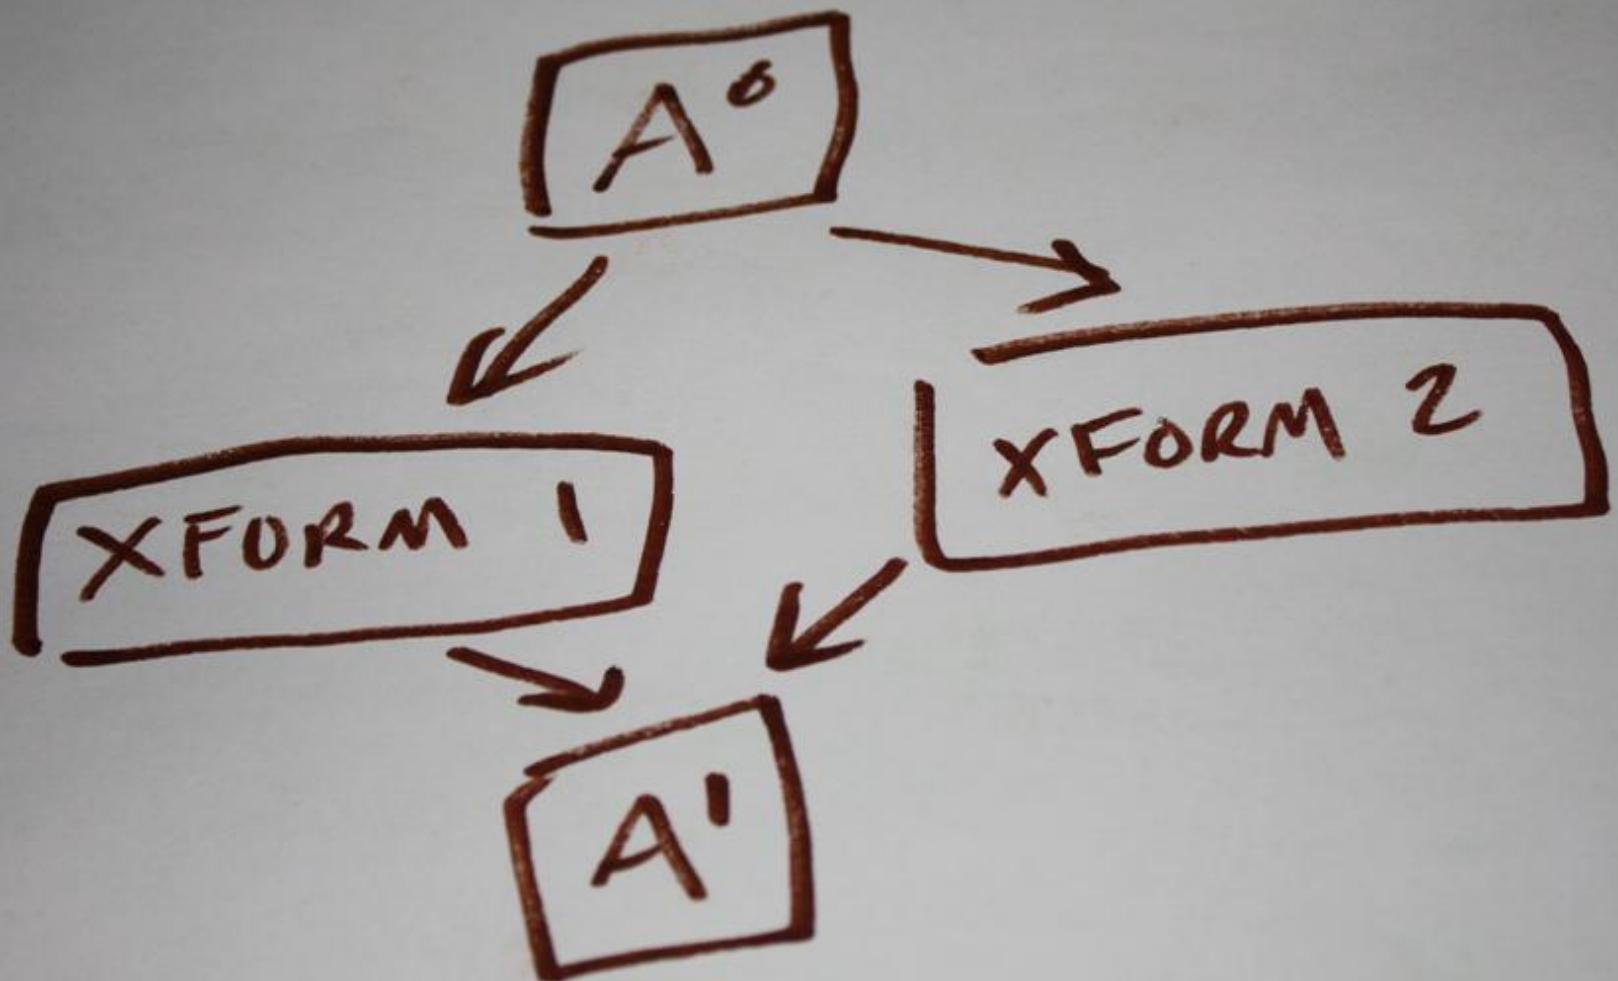
A'

B'

NO PROBLEM  
CAN  
POSSIBLY  
BE  
CAUSED.

A'

B'

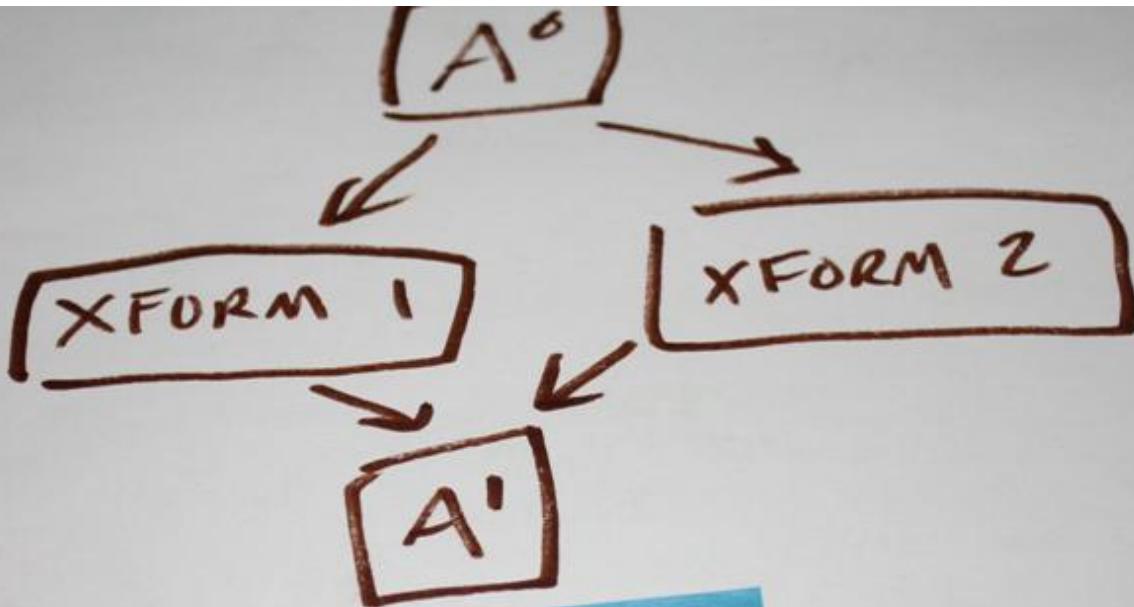


$A^o$

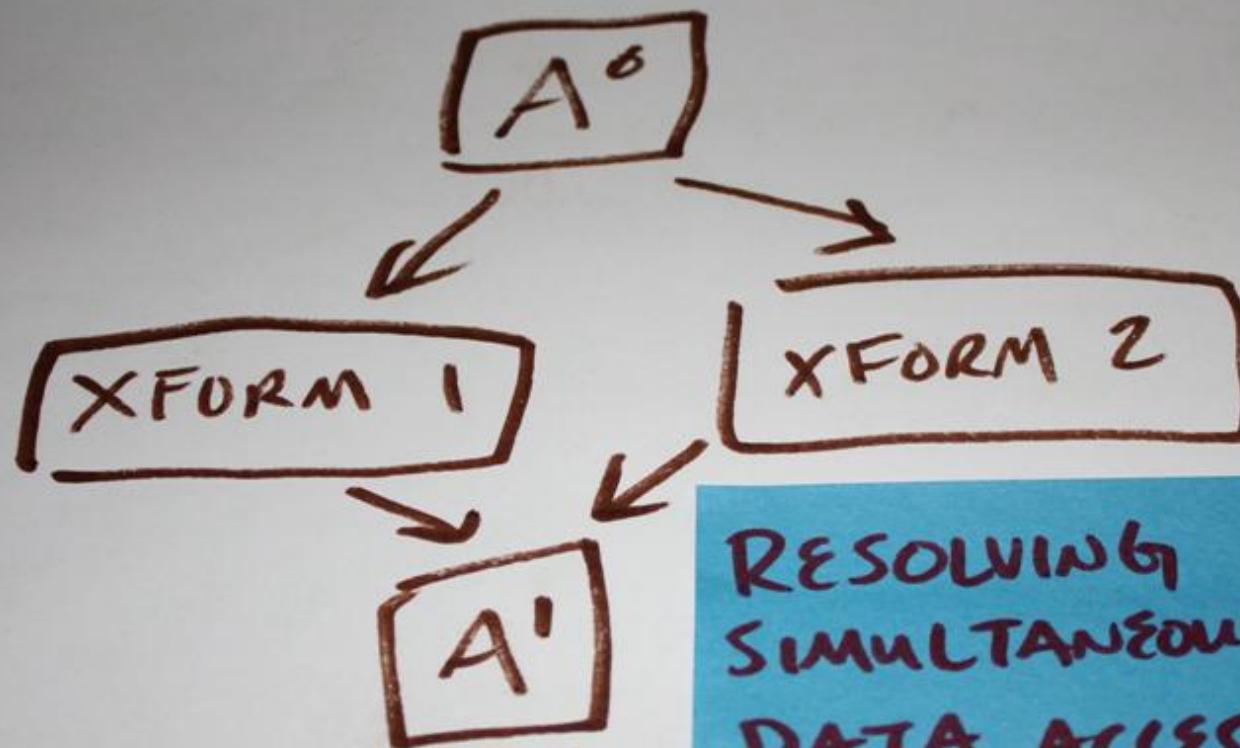
XFORM 1

XFORM 2

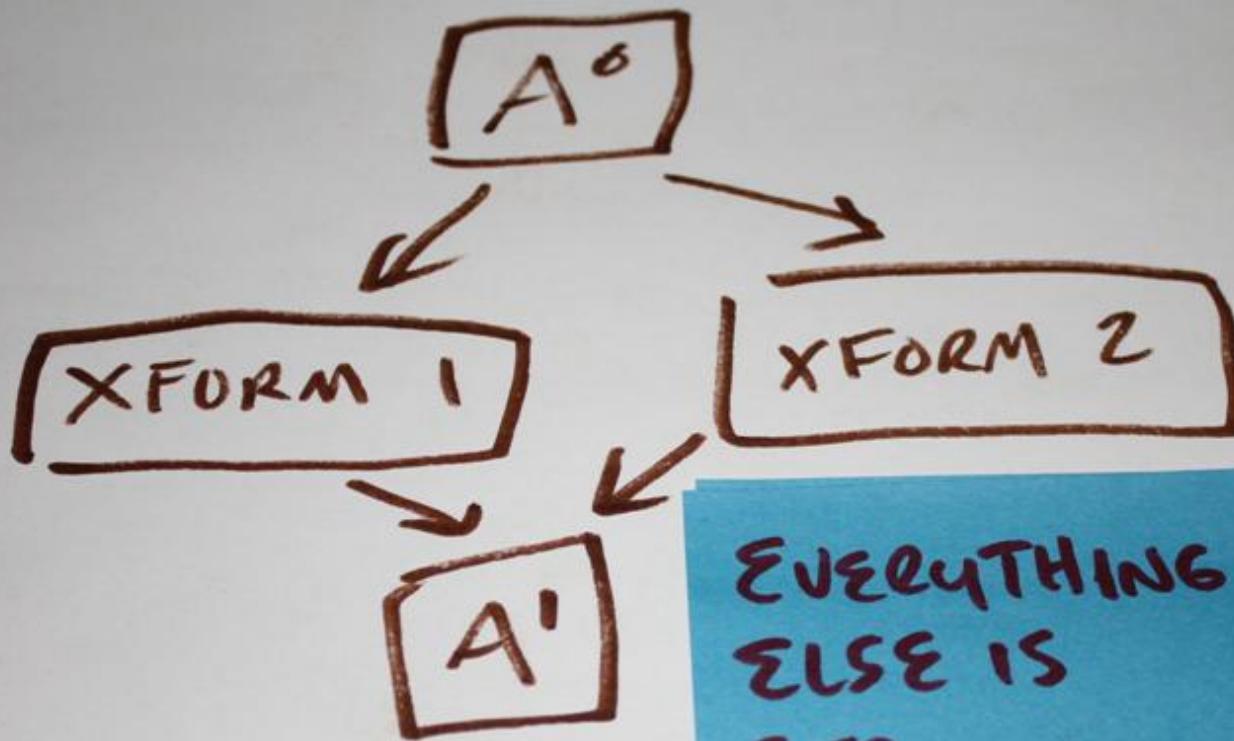
SIMULTANEOUS  
READS FROM  
SAME DATA



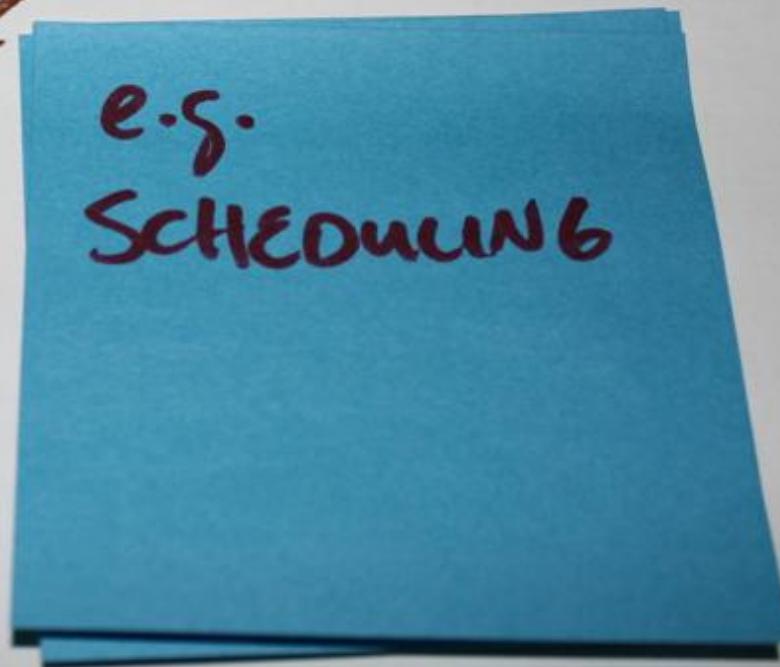
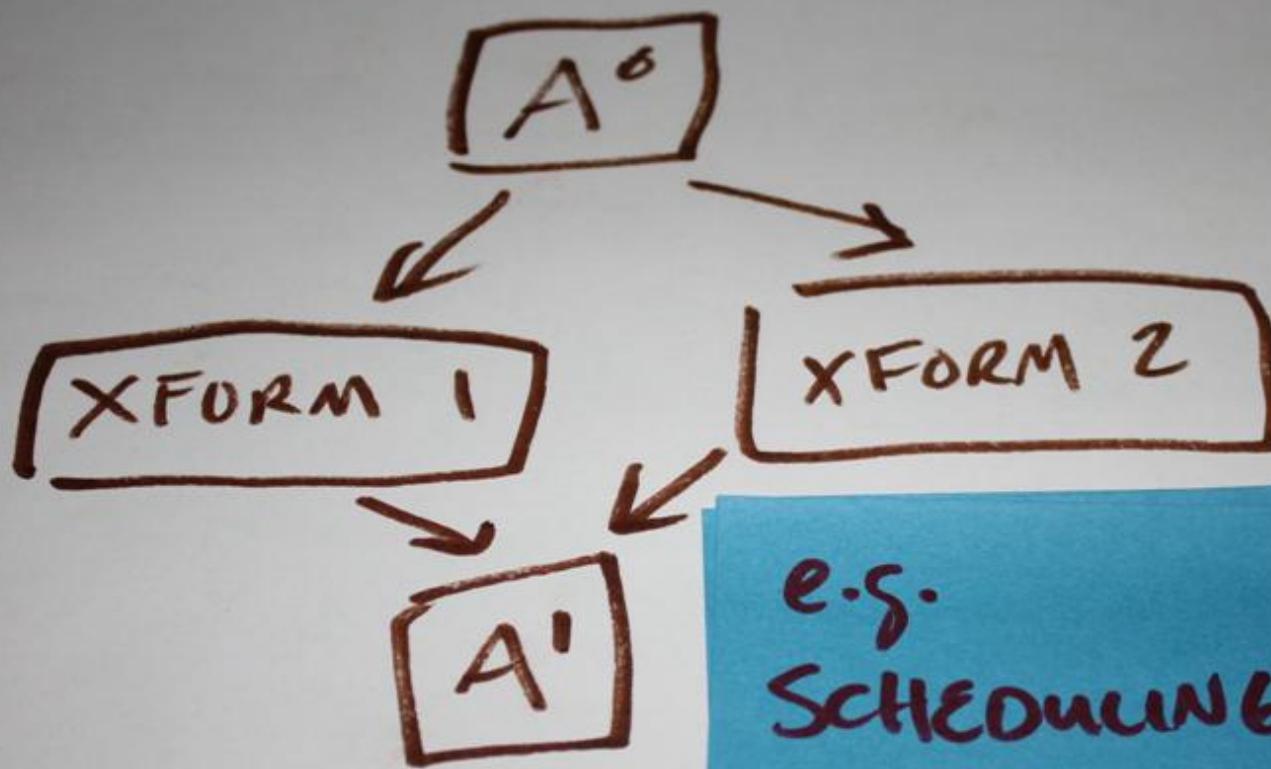
↑  
SIMULTANEOUS  
WRITES TO  
SAME DATA.



RESOLVING  
SIMULTANEOUS  
DATA ACCESS  
IS THE  
CONCURRENCY  
PROBLEM.



EVERYTHING  
ELSE IS  
STD.  
RESOURCE  
MANAGEMENT.



## FIRST - VARIOUS PARALLELISM

- INSTRUCTION LEVEL
- MULTI - THREADING
- MULTI - CORE, SHARED MEM
- MULTI - CORE, INDEPENDENT MEM
- MULTI - MACHINES

WHAT DO ALL OF  
THESE HAVE  
IN COMMON?

SIMULTANEOUS  
XFORM OF  
SAME DATA FILE

CONCURRENCY IS  
AN ATTRIBUTE OF  
AN OPERATION.

CONCURRENCY IS  
AN ATTRIBUTE OF  
AN OPERATION.

WHICH  
OPERATION?

WHAT ABOUT USING  
TIMESTAMPS TO CONTROL  
ORDER?

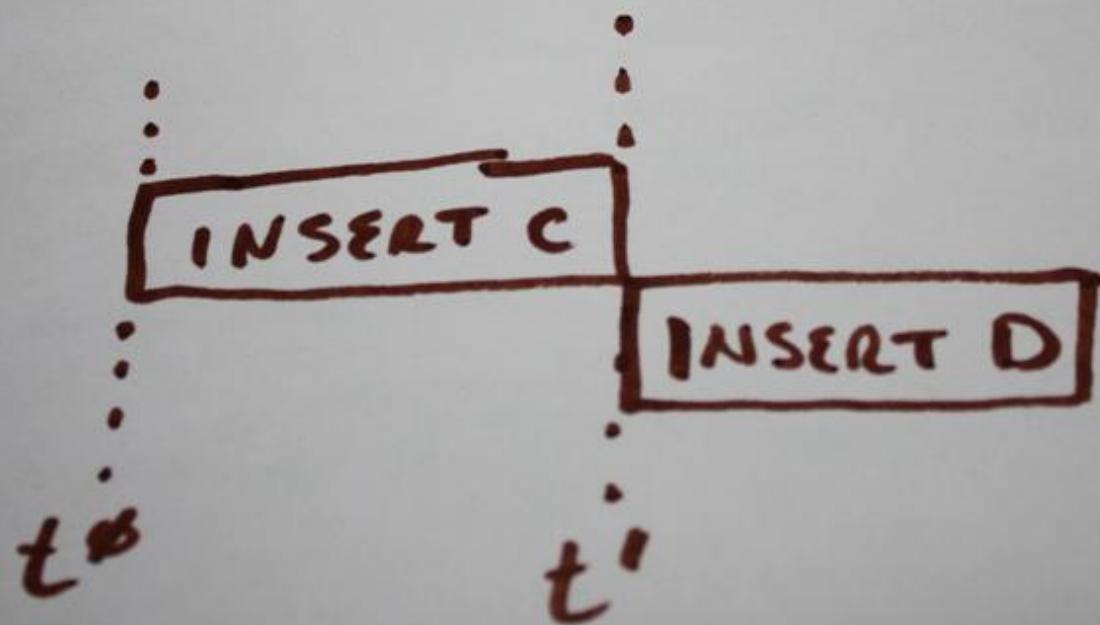
WHAT ABOUT USING  
TIMESTAMPS TO CONTROL  
ORDER?

# I NEED  
Perfectly  
Sync'd  
clocks...

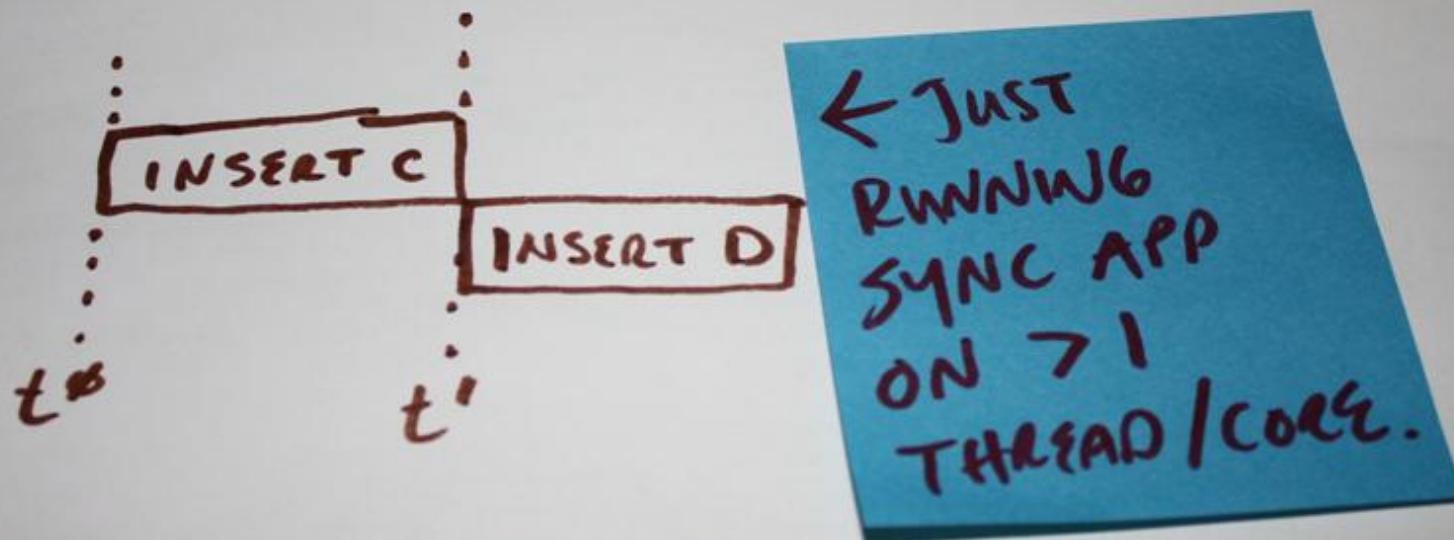
WHAT ABOUT USING  
TIMESTAMPS TO CONTROL  
ORDER ?

...  
which are  
infinitely  
accurate

#2 THIS IS NOT  
CONCURRENCY!



#2 THIS IS NOT CONCURRENCY!



# #2 THIS IS NOT CONCURRENCY!

However. . .

← JUST  
RUNNING  
SYNC

# CONCURRENCY:

Cpus  
aren't  
fully  
concurrent  
either

t<sup>0</sup>

t<sup>1</sup>

ALL THE CONCURRENCY!

There is  
a  
clock

16



# CONCURRENCY

There are  
explicit  
ordering  
rules

$t_0$

$t$

D

Can't be  
more  
concurrent  
than H/W  
allows.

WHAT IS THE  
"ASSEMBLY" OF  
CONCURRENCY?

ASSEMBLY IS THE  
ASSEMBLY OF  
CONCURRENCY .

ASSEMBLY IS THE  
ASSEMBLY OF  
CONCURRENCY.

CRAZY,  
RIGHT?

Concurrent

~ How does  
the H/w  
work?

ATOMIC  
ATOMIC

READ  
WRITE

How do they work  
on the target H/w?

Concurrency is FIRST  
ABOUT RESOLVING DATA  
SYNCHRONIZATION  
CONFLICTS.

Concurrency is FIRST  
ABOUT RESOLVING DATA  
SYNCHRONIZATION  
CONFLICTS.

DEFINE  
DATA IN  
CONTEXT

Concurrency is FIRST  
ABOUT RESOLVING DATA  
SYNCHRONIZATION  
CONFLICTS.

MINIMIZE  
CONFLICTS

Concurrency is FIRST  
ABOUT RESOLVING DATA  
SYNCHRONIZATION  
CONFLICTS.

CONFLICTS =  
SEQUENTIAL  
DATA  
DEPENDENCY

RONIZATION

UCTS.

SEQUENTIAL  
=  
NOT  
CONCURRENT

NEED TO ANSWER BASIC  
QUESTIONS ABOUT THE  
DATA.

WHO      WHAT      WHEN

WHERE      WHY      HOW

NEED TO ANSWER BASIC  
QUESTIONS ABOUT THE  
DATA.

WHO  
CAN  
READ/  
WRITE?

WHAT WHEN  
WHERE WHY HOW

0 TO ANSWER  
QUESTIONS ABOUT THE  
DATA.

DATA  
CAN THE  
DATA BE  
READ OR  
WRITTEN?

# QUESTION DATA

WHO  
CAN  
READ  
WRITE

WHAT DATA  
REALM  
NEEDS TO  
BE R/W?

WHEN  
CAN  
DATA  
BE  
WRITTEN

# QUESTIONS ABOUT DATA.

WHO  
CAN  
READ  
WRITE

HOW IS  
THE DATA  
STORED?

WHEN  
IN  
DATA  
RE  
WE

NEED TO ASK  
QUESTIONS ABOUT THE  
DATA.

WHO CAN READ  
WRITING

FOR HOW LONG IS  
THE DATA  
ACCESSED? IN THE  
DATA AND

DATA IS  
WHAT ARE  
THE  
LATENCIES  
CONSTRAINTS. QD?

DATA

QUESTIONS  
DATA :  
WHAT & WHAT ARE  
THE  
LARGEST THE OUTPUT IN THE  
CONSTRAINTS? TA  
AND  
WHAT

QUESTIONS  
DATA:

WHAT & WHY DO  
YOU NEED  
TO READ  
CONSIDER THE DATA?

CONCURRENT INSERT OP:

WHAT DOES IT  
MEAN IN CONTEXT?

BUT...

- INSERT (C) AFTER (A)
- INSERT (D) AFTER (A)

HAS NO WELL-DEFINED  
MEANING IN A GENERAL  
CONCURRENT SYSTEM!

BUT...

- INSERT (C) AFTER (A)
- INSERT (D) AFTER (A)

THERE IS  
NO "NOW"

IN

CONCURRENCY!

ED

ERAL

!

HAS NO MEANING

BUT...

- INSERT (C) AFTER (A)
- INSERT (D) AFTER (A)

w/out  
"now",  
THERE'S NO  
BEFORE/AFTER

HAS NO MEANING  
EDERAL  
DEPENDENT SYSTEM!

BUT...

- INSERT (C) AFTER (A)
- INSERT (D) AFTER (A)

TIME IS THE  
EXTRA DIM  
THAT MUST  
BE DEFINED.

HAS NO MEANING  
IN CURRENT SYSTEM!

BUT...

- INSERT (C) AFTER (A)
- INSERT (D) AFTER (A)

How IS IT  
HANDLED?

HAS NO MEANING  
ED  
ERAL  
!  
SYSTEM!

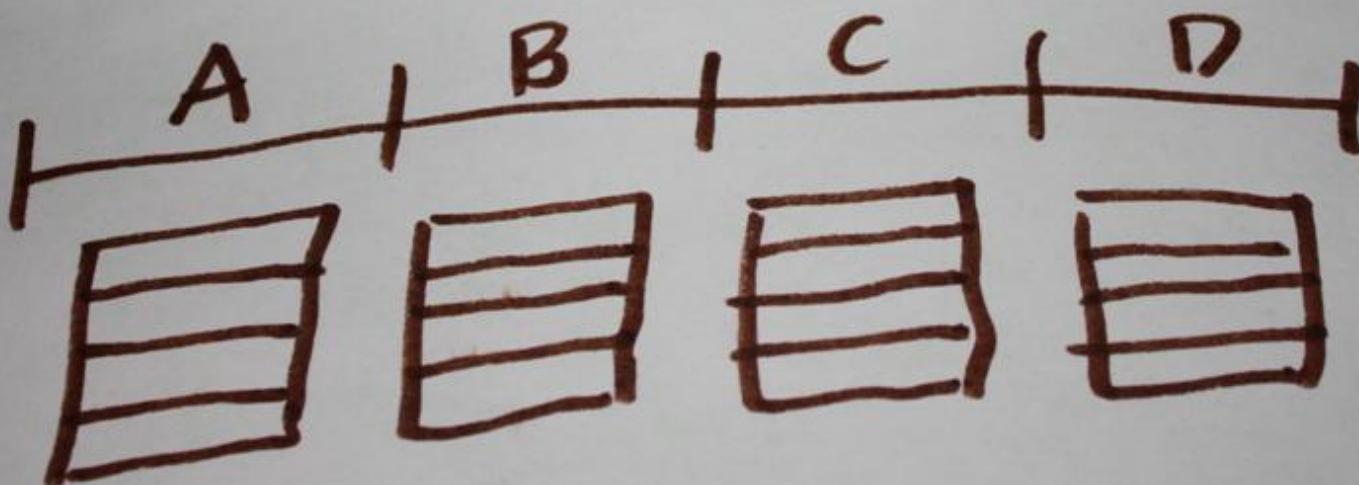
- INSERT (C) AFTER (A)
- INSERT (D) AFTER (A)

WHAT DOES  
IT MEAN?

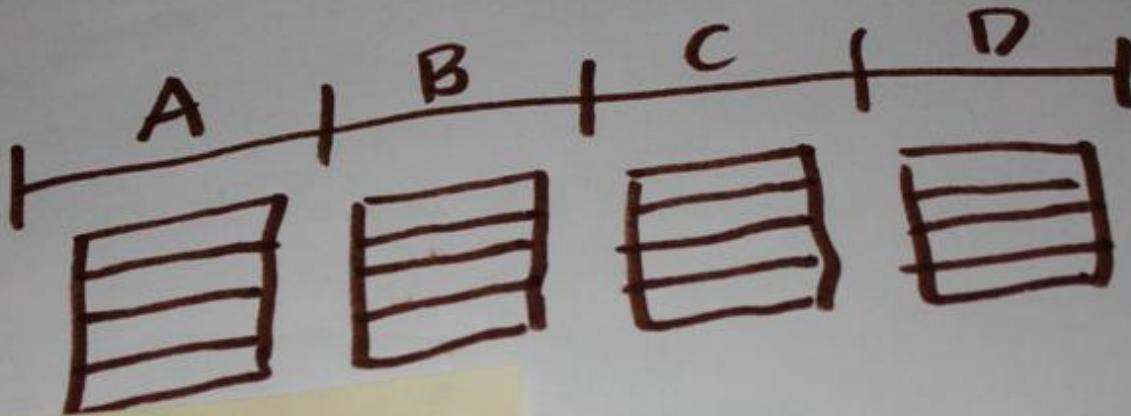
INSERT (D) AFTER (A)

WHAT ARE  
THE LIMITS  
& RANGE?

WHAT IS THE ACCURACY /  
GRANULARITY OF GLOBAL  
ORDER VALUES?

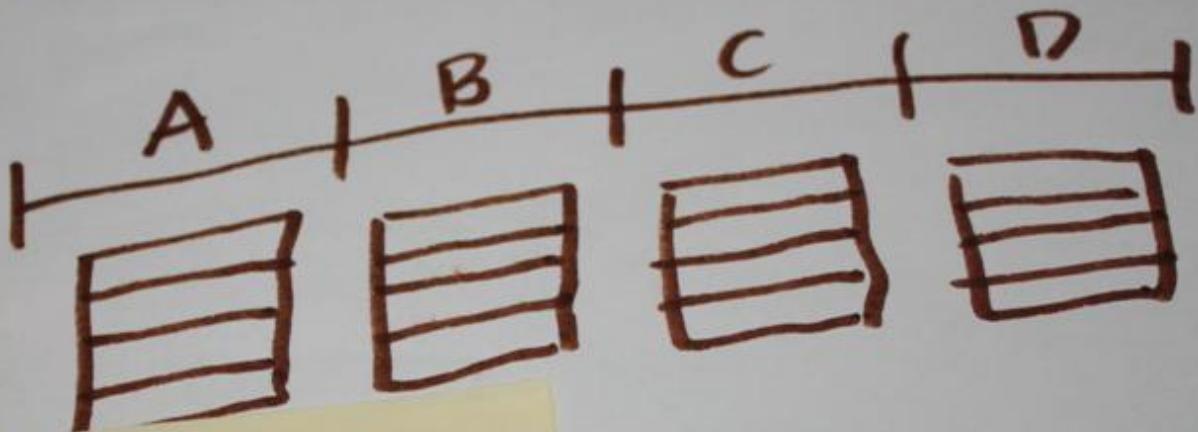


GRANULARITY  
ORDER VALUES?



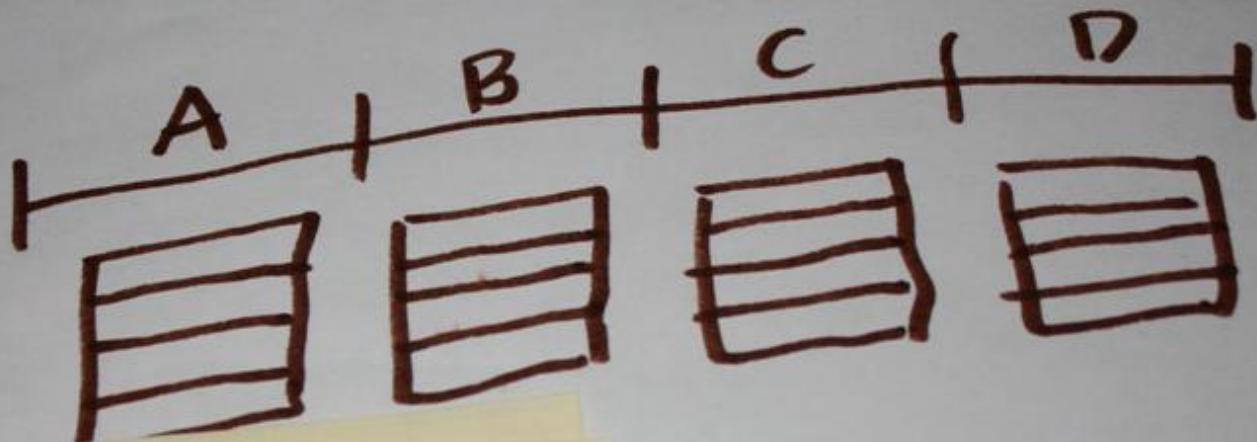
Maybe  
"Insert A"  
means in  
(A) bucket.

order values.



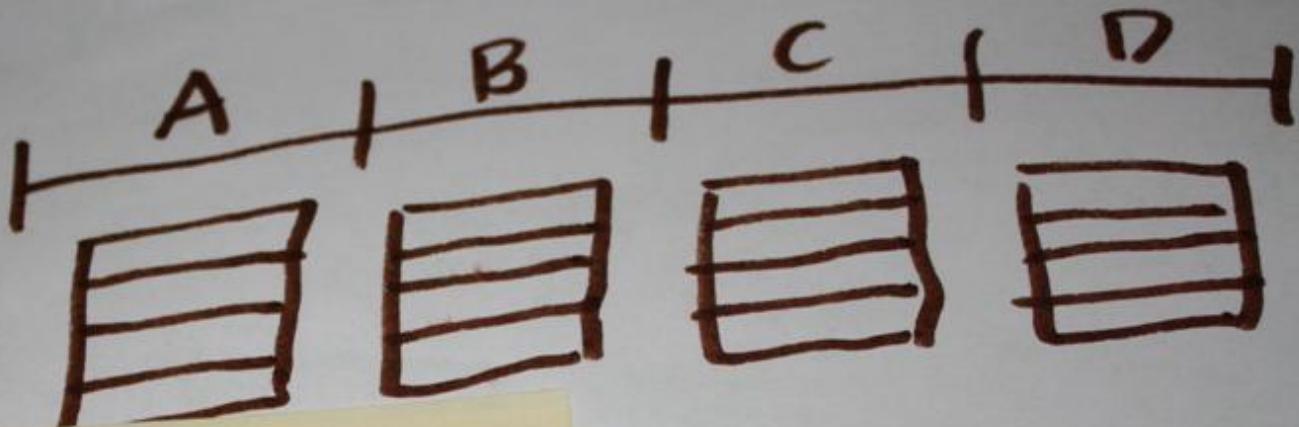
i.e.  
within some  
acceptable  
range.

order



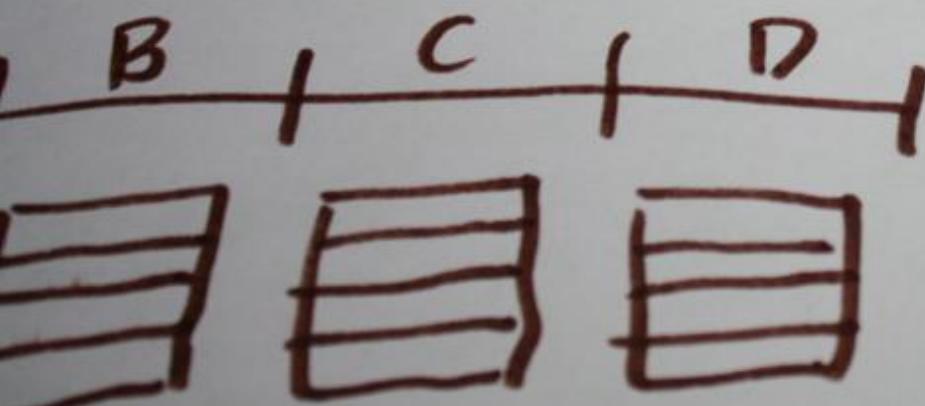
order in  
brackets is  
not  
important

ORDER



but global  
order is  
well-defined.

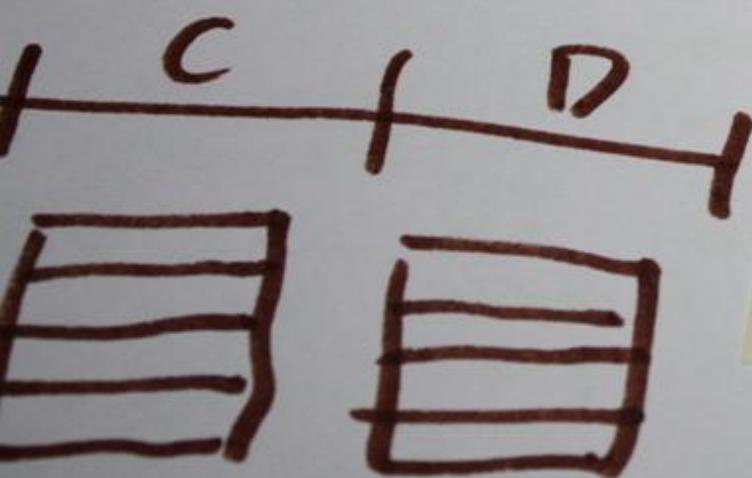
IS THE ACCURACY /  
UNIQUENESS OF GLOBAL  
VALUES?



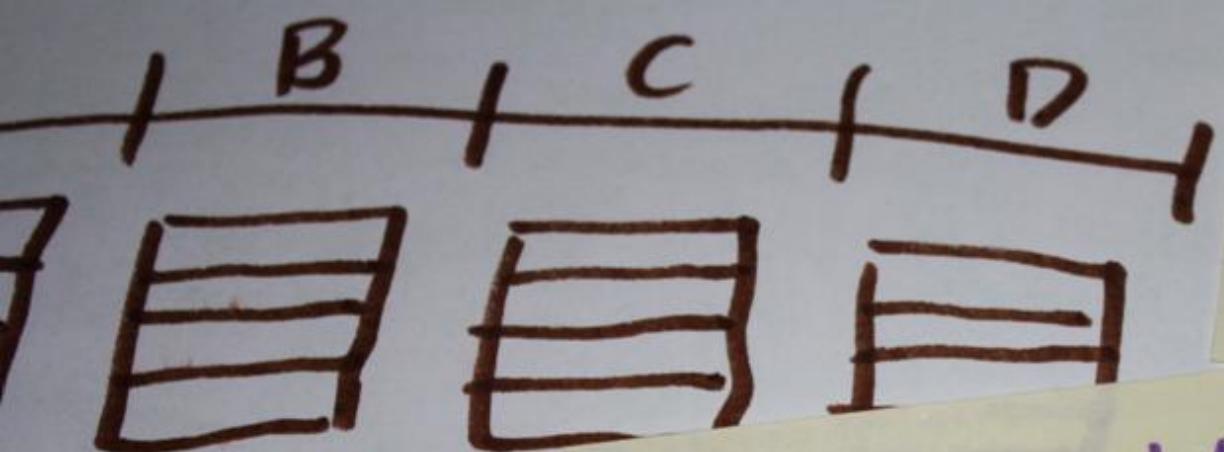
This is an example of an ordering rule.

W  
n  
d

THE ACCURACY /  
PRECISION OF GLOBAL  
VALUES?



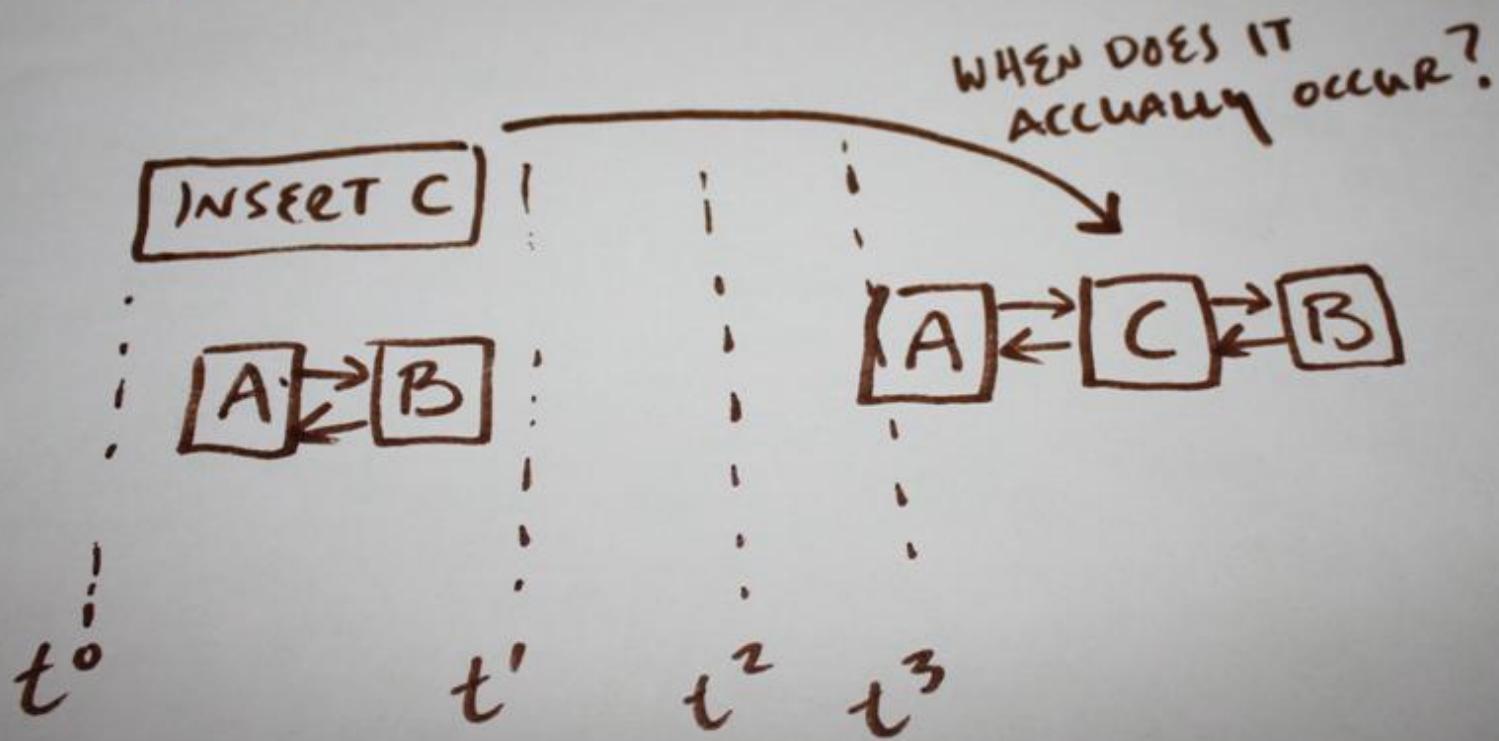
which is  
necessary to  
define in  
concurrent  
problems.



define in  
concurrent  
problems.

How would  
this make  
the insert  
OP simpler?

WHAT ARE THE LATENCY  
REQUIREMENTS OF THE  
INSERT OP?



## LESSON

DOUMLY-LINKED LIST IS  
SEQUENTIAL DATA STRUCT.

PRESUMES:

- ZERO LATENCY
- GUARANTEED ORDER  
(SEQUENTIAL)

# CONCURRENT DATA STRUCTURES DEFINED

By:

- EXPLICIT LATENCY  
REQUIREMENTS
- EXPLICIT ORDERING  
RULES

WHAT ARE SOME  
CONTEXTS W/IN YOU  
WOULD USE DOUBLY-  
LINKED LIST?

WHAT ARE SOME  
CONTEXTS W/IN YOU  
WOULD USE DOUBLY-  
LINKED LIST?

I.e.  
DON'T TRY  
TO FIT THE  
"SOLUTION"  
TO THE  
PROBLEM.

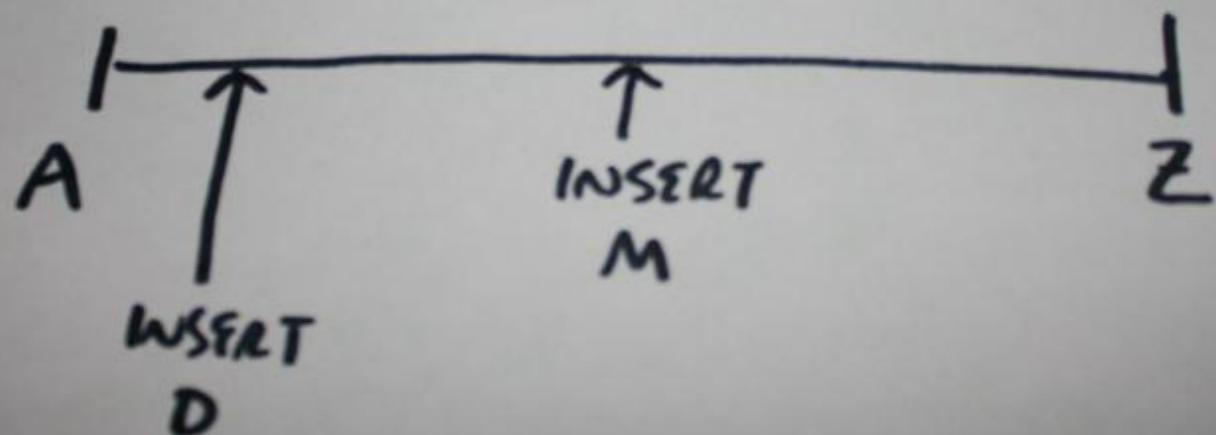
WHAT ARE SOME  
CONTEXTS W/IN YOU  
WOULD USE DOUBLY-  
LINKED LIST?

i.e. ... IT TRY

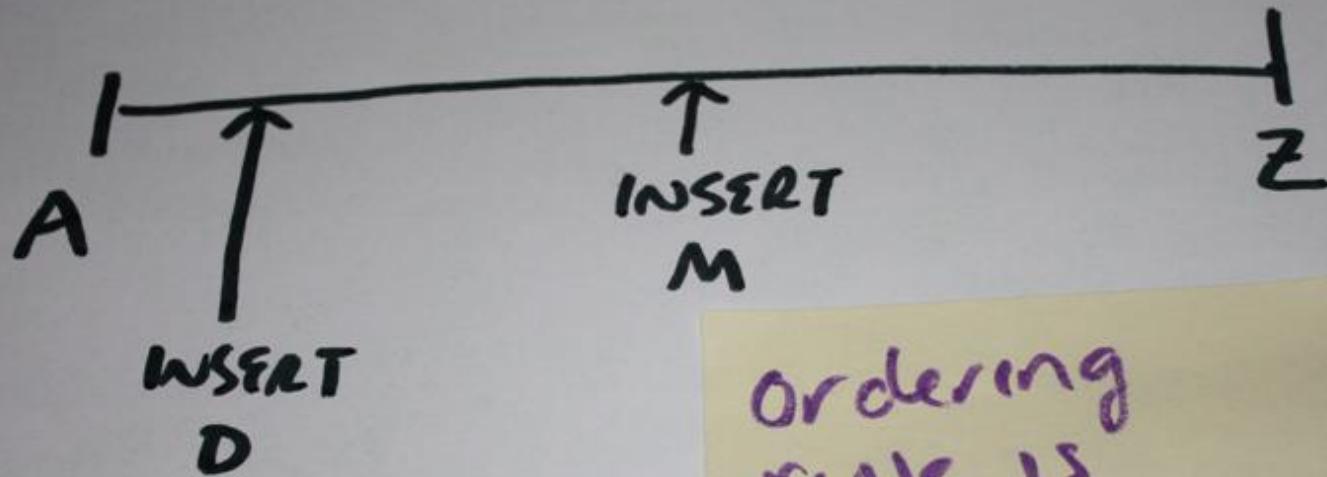
JUST  
SOLVE THE  
PROBLEM.

e.g. Insert  
Sort

INSERT SORT MEANS  
INSERT AT POSITION IN  
GLOBAL ORDER



# GLOBAL ORDER



Ordering rule is global  
Sorting /  
compose func.

AT POSITION IN  
ORDER

---

↑  
INSERT  
M

But what  
are the  
latency  
rules?

ordering

POSITION

ER

When does  
the caller  
need the  
results?

↑  
INSERT

When do  
other  
processes  
need  
results?

# GLOBAL ORDER

Do you think dbl  
linked list  
will be  
right data?

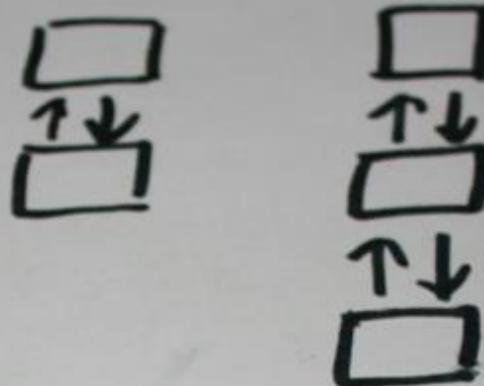


INSERT  
D

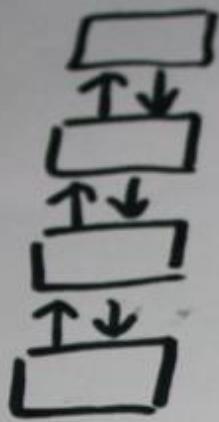
e.g. Resource Mgmt  
w/ variable  
life times



LIMITED  
RESOURCE  
LIST

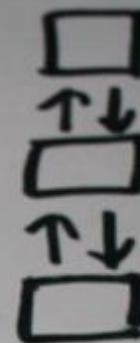
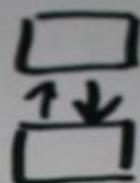


- SPARSE LISTS OF ALLOCATED NODES
- VARIABLE LIFETIMES
- DIFF "OWNERS", DIFF LISTS

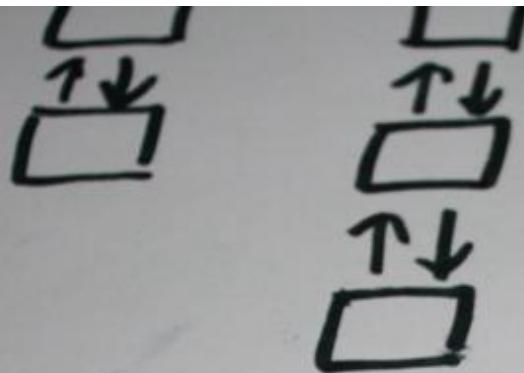


LIMITED  
RESOURCE  
LIST

↑  
Start w/  
LIST OF  
RESOURCES



- SPARSE LISTS OF ALLOCATED NODES
- VARIABLE LIFETIMES
- DIFF "OWNERS", DIFF LISTS



[ ]

- SPARSE LISTS OF ALLOCATED NODES
- VARIABLE LIFETIMES
- DIFF "OWNERS",  
DIFF LISTS

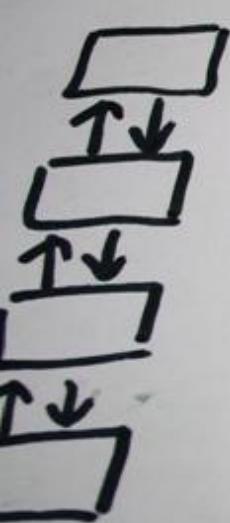
< AS They're alloc'd, add to new "owner" lists.



L

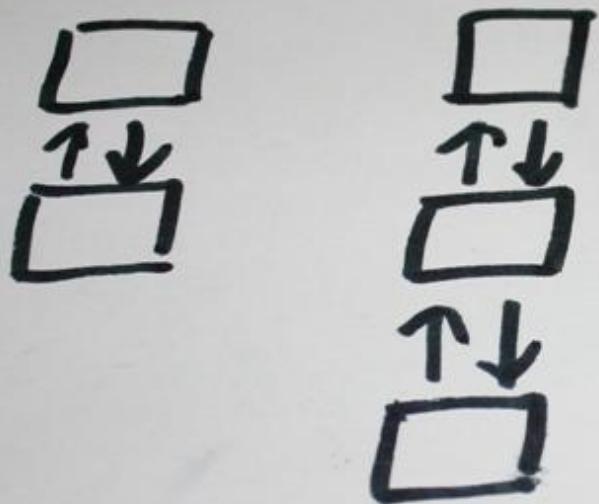
- SPARSE LISTS OF ALLOCATED NODES
- VARIABLE LIFETIMES
- DIFF "OWNERS",  
DIFF LISTS

← When  
"deleted"  
remove from  
list, return  
to res. list.

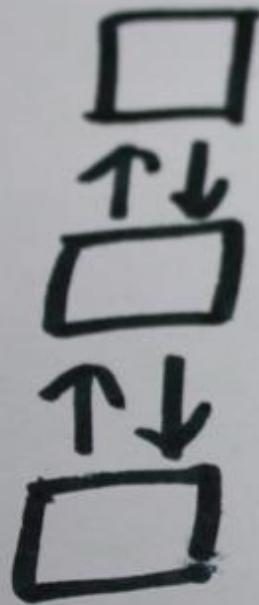
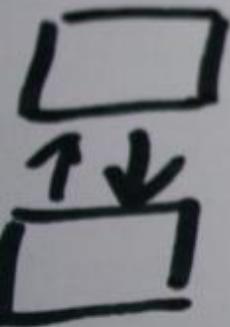


← APPEND  
TO END.  
ORDER NOT  
IMPORTANT

- SPARSE LISTS OF ALLOCATED NODES
- VARIABLE LIFETIMES



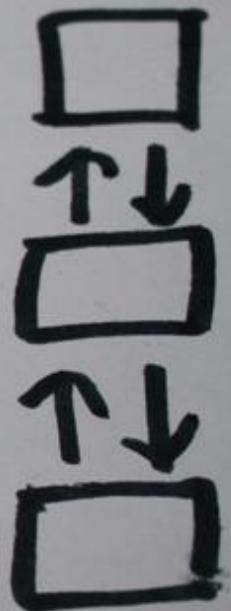
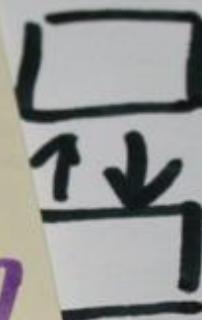
- SPARSE LISTS OF  
ALLOCATED NODES



Note, implies  
insert to  
middle is  
not needed.

process lists of

Again,  
what are  
the latency  
requirements?



LIMITED

- SPARSE LISTS

DATA FOR DIFFERENT  
ORDERING RULES WILL  
BE DIFFERENT.

DATA for different  
latency rules will  
be different.

So...  
why would we  
expect data structs  
for sequential rules  
& concurrent rules to  
be the same?

Begin at  
the beginning

GOOD CONCURRENCY  
DOESN'T ADD UNNEEDED  
SYNC POINTS.

GOOD CONCURRENCY  
DOESN'T ADD UNNEEDED  
SYNC POINTS.

BUT SHOULD  
DESIGN  
AROUND  
PRE-EXISTING  
ONES.

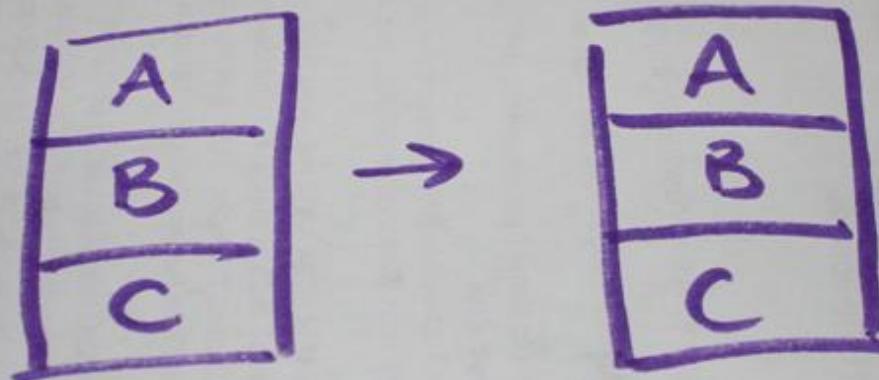
GOOD CONCURRENCY  
DOESN'T ADD UNNEEDED  
SYNC POINTS.

OR USE  
CHEAPEST  
ONES  
AVAILABLE

Concurrency starts with  
atomic transaction



IN-ORDER STORES



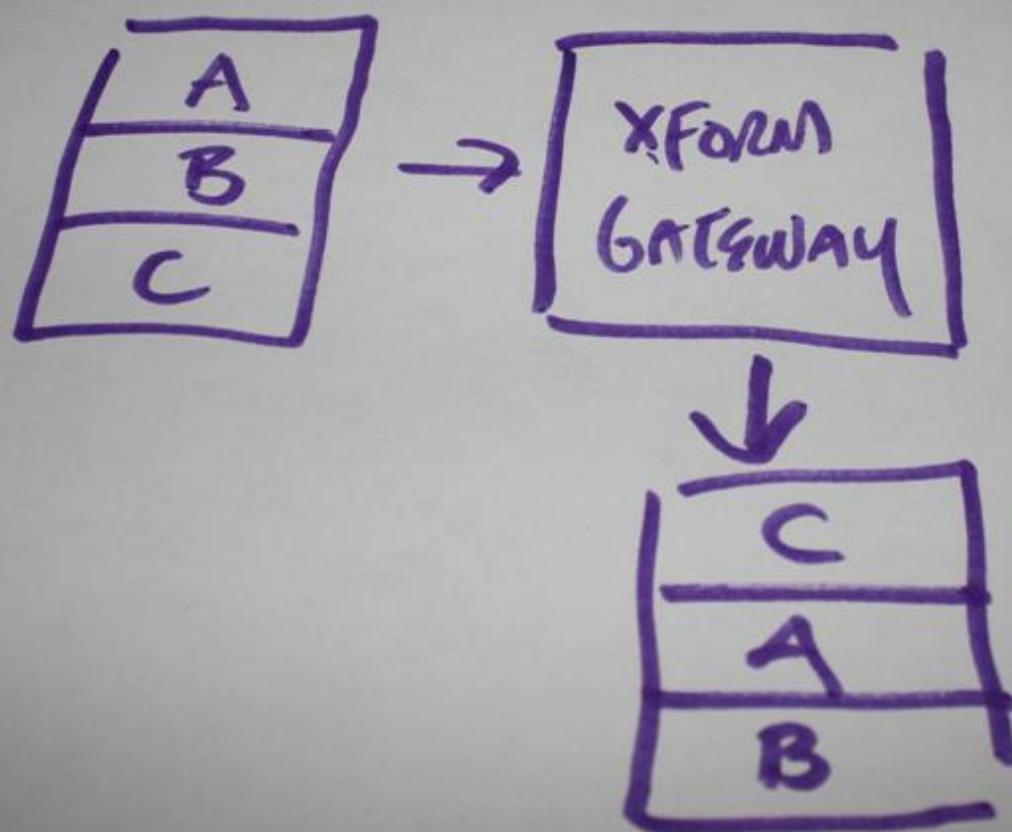
FIFO, NO CHANGE

IN-ORDER  
PROCESSORS HAVE  
IMPLICIT SYNC PT  
EACH INSTRUCTION!

IN-ORDER  
PROCESSORS HAVE  
IMPLICIT SYNC PT  
EACH INSTRUCTION!

↑  
HERE USE  
PRE-EXISTING  
SYNC PTS.

## OUT-OF-ORDER STORES



OUT-OF-ORDER  
PROCESSORS PROVIDE  
ORDERING PRIMITIVES  
(e.g. FENCE)

OUT-OF-ORDER  
PROCESSORS PROVIDE  
ORDERING PRIMITIVES  
(e.g. FENCE)

BUT, MAY  
BE EVEN  
CHEAPER  
OPTION!

PROCESSORS PROVIDING  
MEMORY ORDERING PRIMITIVES  
(e.g. FENCE)

IF CAN SUPPORT HIGHER LATENCY,  
ADD...

ESSORS  
ING PRIMITIVES  
(e.g. FENCE)

HIGH-  
LEVEL SYNC  
PT ALREADY  
EXISTS.

ALSO NOTE

LANGUAGE/COMPILER

Compiler / optimizer re-orders  
instructions by definition!

Prog. must force  
order!

NEED TO ANSWER :

- HOW WILL DATA BE TRANSFORMED?
- WHAT ARE THE CONSTRAINTS?

NEED TO ANSWER

- HOW WILL DATA BE TRANSFORMED?

- WITH  
CON

i.e.  
operations

↑

- WHAT ARE THE  
CONSTRAINTS?

↑  
TO DATA

- WHAT ARE THE  
CONSTRAINTS?



TO  
TRANSFORMATION

- WHAT ARE THE  
CONSTRAINTS?



TO  
GREENING

- WHAT ARE THE  
CONSTRAINTS?

↑

TO  
LATENCY

- WHAT ARE THE  
CONSTRAINTS?

↑

TO

GLOBAL  
GUARANTEES

- WHAT ARE THE  
CONSTRAINTS?



TO

LOCAL  
GUARANTEES

THE SECRET  
of OPTIMIZED  
Concurrent design.

IS...

DELAY\*

\* kind of ironic ,right ?

THE SOONER YOU NEED  
SOME DATA, THE SLOWER  
THE SYSTEM WILL BE.

THE SOONER YOU NEED  
SOME DATA , THE SLOWER  
THE SYSTEM WILL BE .

( LATENCY VS. THROUGHPUT )

THAT'S  
IT!  
END OF  
CLASS. ☺

ENCE

IGHF